

Education

Master of Science in Data Science

University of Colorado Boulder, Colorado

GPA: 3.9/4.0

Aug 2024 – May 2026

Bachelor of Technology in Information Technology / Minor in Data Science

Dwarkadas J. Sanghvi College of Engineering, University of Mumbai, India

GPA: 8.93/10.0

Aug 2020 – May 2024

Work Experience

Temporary Post Graduate Assistant

University of Colorado Boulder

Jul 2026 – Present

Remote, USA

- Designed two distinct LLM persona specifications (GPT-4o) for real-time voice-based role-play, defining behavioral rules that vary response depth and information disclosure based on live conversational input quality.
- Engineered a speech-to-text → LLM → text-to-speech pipeline for natural, low-latency spoken conversation, validating stage-specific prompts through iterative live testing.
- Architected a 5-stage sales simulation platform with automated LLM-based grading, translating qualitative rubric criteria into structured, machine-scorable evaluation prompts across 28 distinct behavioral checks.
- Built a config-driven synthetic dataset generator with a multi-axis comparison engine and implemented Row-Level Security across a multi-table PostgreSQL (Supabase) schema, directing implementation as technical lead.

AI Specialist Intern

Projxon

Aug 2025 – May 2026

Remote, USA

- Built an automated ETL pipeline integrating LinkedIn and Windsor.ai APIs into Google Sheets → BigQuery → Looker Studio, designing structured schemas and SQL transformations for KPI aggregation.
- Applied Association Rule Mining to analyze content patterns and engagement relationships, building dashboards for audience demographics, engagement trends, and post performance that automated manual reporting.
- Built end-to-end HR automation pipelines in Google Apps Script, automating dynamic document generation with validation logic, role-based templating, and batch processing.
- Structured OrcaAI, an AI-powered business optimization consultant for SMBs using predictive analytics, classification, and recommendation logic, securing 2nd place and \$25K in AWS credits.

AI/ML Engineer Intern

Exaplus Technologies

Dec 2023 – Jul 2024

India

- Developed an automated financial reconciliation system that ingests invoices, bank statements, and spreadsheets from multiple sources, using NLP-based extraction logic and statistical comparison rules to flag anomalous entries.
- Built a K-Means clustering model (Scikit-learn/Pandas) that segments clients by transaction behavior patterns, engineering features from raw transaction records to auto-route clients to report templates.
- Built an automated document-intelligence pipeline using the OpenAI API to extract structured JSON from unstructured client reports and financial summaries, exposed as an internal FastAPI endpoint.

Technical Skills

Languages: Python, R, SQL.

ML & Statistical Modeling: Machine Learning, Deep Learning, Gradient Boosting, Naive Bayes, K-Means Clustering, OLS/Ridge Regression, Association Rule Mining.

Data Analysis & Visualization: Pandas, Tableau, Power BI, DAX, Looker Studio, GGplot.

Databases & Cloud: PostgreSQL, MySQL, BigQuery, SQLite, AWS, GCP.

NLP & Computer Vision: NLP, Computer Vision, OCR.

Responsible AI: SHAP, LIME, ALE, Bias Awareness.

Projects

Asteroids Analysis | PostgreSQL, R, Naive Bayes, GGplot, Liquibase

- Built a Naive Bayes classification pipeline for near-earth and hazardous asteroids, engineering features across 100,000+ records in PostgreSQL.
- Achieved 98% classification accuracy; used Liquibase for database version control and reproducible schema migrations.

Steam Game Insights | Gradient Boosting, CLIP, Association Rule Mining, Tableau

- Analyzed 25,000+ Steam games via API data, applying CLIP embeddings and Gradient Boosting for success prediction.
- Applied Association Rule Mining for genre and pricing relationships; built Tableau dashboards for trend visualization.

FIFA Player Analytics Dashboard | Power BI, DAX

- Built an interactive Power BI dashboard analyzing 17,000+ FIFA players across multiple datasets with advanced DAX measures and dynamic player comparisons.
- Built position-based scatter plots and custom tooltips for player comparison analysis.

Meridian | Scikit-learn, SciPy, FastAPI, OpenAI API

- Implemented geometric adstock decay, Hill saturation functions, and OLS/Ridge regression to decompose revenue into channel contributions and estimate ROI/elasticity, with model diagnostics (R^2 , MAPE).
- Developed a SciPy SLSQP budget optimizer recommending optimal channel allocation under diminishing returns.